# Deep Generative Models Programming Exersice 2

Ron Ferens, 037222825

## Code

Implementation of the VAE.py file:

```
"""NICE model
import torch
import torch.nn as nn
import torch.nn.functional as F
import typing
class Model(nn.Module):
    def __init__(self, latent_dim: int, device: str):
         Initialize a VAE
              latent dim: dimension of embedding
              device: run on cpu or gpu
         # Encoder
         # -----
         super(Model, self). init ()
         self.device = device
         self.latent_dim = latent_dim
         self.encoder = nn.Sequential(
              nn.Conv2d(1, 32, 4, 1, 2), # B, 32, 28, 28
              nn.ReLU(True),
              nn.Conv2d(32, 32, 4, 2, 1), # B, 32, 14, 14
              nn.ReLU(True),
              nn.Conv2d(32, 64, 4, 2, 1), # B, 64, 7, 7
         )
         self.mu = nn.Linear(64 * 7 * 7, latent dim)
         self.log var = nn.Linear(64 * 7 * 7, latent dim)
         # -----
         # Decoder
         self.upsample = nn.Linear(latent dim, 64 * 7 * 7)
         self.decoder = nn.Sequential(
              nn.ConvTranspose2d(64, 32, 4, 2, 1), # B, 64, 14, 14
              nn.ReLU(True),
              nn.ConvTranspose2d(32, 32, 4, 2, 1, 1), # B, 32, 28, 28
              nn.ReLU(True),
              nn.ConvTranspose2d(32, 1, 4, 1, 2), # B, 1, 28, 28
              nn.Sigmoid()
         )
    def sample(self, sample size: int, mu: float = None, log var=None) -> typing.List:
         Sampled images from the model
         :param sample size: Number of samples
         :param mu: z mean, None for prior (init with zeros)
         :param log var: z log-STD, None for prior (init with zeros)
         :return:
         if mu is None:
              mu = torch.zeros((sample size, self.latent dim)).to(self.device)
         if log var is None:
```

```
log var = torch.zeros((sample size, self.latent dim)).to(self.device)
          # Making sure the model is not trained while sampling and no gradients are
accumulated
         samples = []
         with torch.no grad():
              z = self.z sample(mu=mu, log_var=log_var)
              output = self.decoder(self.upsample(z).view(-1, 64, 7, 7))
              for i in range(sample size):
                   samples.append(output[i].squeeze().data.cpu().numpy())
         return samples
    def z sample(self, mu: torch.Tensor, log var: torch.Tensor) -> torch.Tensor:
         Applying the reparameterization trick to generate Z
         :param mu: Z mean
         :param log var: z log-STD
         :return: samples of Z
         # Getting the STD from log-variance input
         std = torch.exp(0.5 * log var)
         # Sampling epsilon from Normal (0, 1)
         eps = torch.rand like(std).to(self.device)
         # Applying the reparameterization trick for Gaussian
         z = mu + eps * std
         return z
    @staticmethod
    def loss(x: torch.Tensor, recon: torch.Tensor, mu: float, log var: float):
         Calculating the model's loss composed of BCE and KL elements
         :param x: input data
         :param recon: reconstructed data
         :param mu: Z mean
         :param log var: Z log-STD
         :return:
          # Calculating the binary cross entropy loss
         bce loss = F.binary cross entropy(recon, x, reduction='sum')
         # Calculating the KL Divergence loss
         kl loss = torch.sum(1 + log var - torch.exp(log var) - torch.pow(mu, 2)) / 2.0
          # Calculating the ELBO(p, q, x)
         elbo = bce loss - kl loss
         return elbo
    def forward(self, x: torch.Tensor) -> typing.Tuple[torch.Tensor, torch.Tensor, torch.Tensor]:
         Running the model's forward pass
         :param x: input data
         # Encoder forward pass
         x = self.encoder(x)
         x = x.reshape(-1, 64 * 7 * 7)
         # Generating the mean and variance values
         mu = self.mu(x)
         log_var = self.log_var(x)
         # Decoder forward pass
         z = self.z sample(mu=mu, log var=log var)
         output = self.decoder(self.upsample(z).reshape(-1, 64, 7, 7))
```

#### Implementation of the train.py file:

```
"""Training procedure for NICE.
111111
import argparse
import torch, torchvision
from torchvision import transforms
import numpy as np
from VAE import Model
import matplotlib.pyplot as plt
import torch.nn as nn
PLOT DIM = 5
def train(vae: nn.Module, trainloader: torch.utils.data.DataLoader, optimizer: torch.optim) ->
float:
     # set to training mode
    vae.train()
     # Setting an accumulated epoch loss
    running loss = 0
    num iterations = 0
     # Going over the training dataset (single epoch)
    for inputs, in trainloader:
          # Loading the batch input
         inputs = inputs.to(vae.device)
         # Model forward pass
         model res, mu, \log var = vae(inputs)
          # Calculating the loss criteria
         optimizer.zero grad()
         loss = vae.loss(x=inputs, recon=model res, mu=mu, log var=log var)
         running loss += loss
         num iterations += 1
         # Backprop and optimization
         loss.backward()
         optimizer.step()
    epoch loss = running loss / num iterations
    return epoch loss.item()
def test(vae: nn.Module, testloader: torch.utils.data.DataLoader) -> float:
    vae.eval() # set to inference mode
    with torch.no grad():
         running loss = 0
         for batch_idx, (inputs, _) in enumerate(testloader):
              # Loading the batch input
              inputs = inputs.to(vae.device)
              # Model forward pass
              model res, mu, log var = vae(inputs)
              # Calculating the loss criteria
              loss = vae.loss(inputs, model res, mu=mu, log var=log var)
              running loss += loss.item()
    test loss = running loss / (batch idx + 1)
    return test loss
```

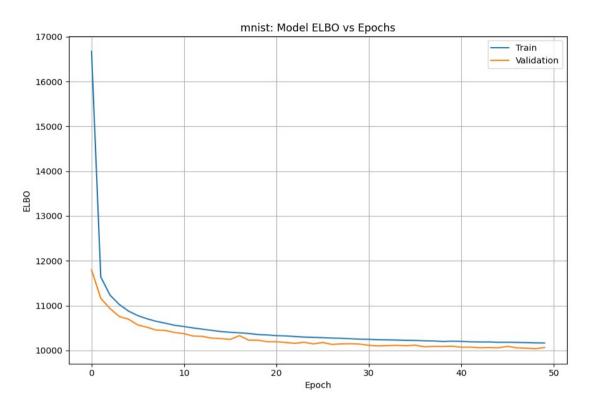
```
def main(args):
    device = torch.device("cuda:0" if torch.cuda.is available() else "cpu")
    transform = transforms.Compose([
         transforms.ToTensor(),
         transforms.Lambda(lambda x: x + torch.zeros like(x).uniform (0., 1. / 256.)), #
dequantization
         transforms.Normalize((0.,), (257. / 256.,)), # rescales to [0,1]
    ])
    if args.dataset == 'mnist':
         trainset = torchvision.datasets.MNIST(root='./data/MNIST',
                                                       train=True, download=True,
transform=transform)
         trainloader = torch.utils.data.DataLoader(trainset,
                                                            batch size=args.batch size,
shuffle=True, num_workers=2)
         testset = torchvision.datasets.MNIST(root='./data/MNIST',
                                                      train=False, download=True,
transform=transform)
         testloader = torch.utils.data.DataLoader(testset.
                                                           batch size=args.batch size,
shuffle=False, num workers=2)
    elif args.dataset == 'fashion-mnist':
         trainset = torchvision.datasets.FashionMNIST(root='~/torch/data/FashionMNIST',
                                                               train=True, download=True,
transform=transform)
         trainloader = torch.utils.data.DataLoader(trainset,
                                                            batch size=args.batch size,
shuffle=True, num workers=2)
         testset = torchvision.datasets.FashionMNIST(root='./data/FashionMNIST',
                                                              train=False, download=True,
transform=transform)
         testloader = torch.utils.data.DataLoader(testset,
                                                           batch size=args.batch size,
shuffle=False, num workers=2)
    else:
         raise ValueError('Dataset not implemented')
    # Creating an instance of the VEA model
    vae = Model(latent dim=args.latent dim, device=device).to(device)
    # Defining an Adam optimizer
    optimizer = torch.optim.Adam(vae.parameters(), lr=args.lr)
    # Resetting the ELBO arrays for train and validation
    elbo train = []
    elbo val = []
    for epoch in range(args.epochs):
         # Running a single training epoch
         epoch loss = train(vae=vae, trainloader=trainloader, optimizer=optimizer)
         elbo train.append(epoch loss)
         # Running model validation
         val loss = test(vae=vae, testloader=testloader)
         elbo val.append(val loss)
         print(f'Epoch {epoch}: Train Loss={epoch loss}, Validation Loss={val loss}')
    # Plotting the model's train and validation ELBO
    plt.figure()
    plt.plot(elbo train, label='Train')
    plt.plot(elbo val, label='Validation')
    plt.title(f"{args.dataset}: Model ELBO vs Epochs")
```

```
plt.xlabel("Epoch")
    plt.ylabel("ELBO")
     plt.legend()
     plt.grid(True)
     plt.show()
     # Saving samples generated by the trained model
     samples = vae.sample(sample size=args.sample size, mu=None, log var=None)
     fig, ax = plt.subplots(nrows=np.ceil(args.sample size / PLOT DIM).astype(np.uint8),
                                ncols=PLOT DIM)
    for i in range(args.sample size):
         ax[(i // PLOT DIM), (i % PLOT DIM)].imshow(samples[i])
    [axi.set axis off() for axi in ax.ravel()]
     fig.suptitle(f"{args.dataset}: Trained Model - Output Samples")
     plt.show()
if __name__ == '__main__':
    parser = argparse. Argument Parser('')
    parser.add argument('--dataset',
                             help='dataset to be modeled.',
                             type=str,
                             default='fashion-mnist')
    parser.add argument('--batch size',
                             help='number of images in a mini-batch.',
                             type=int,
                             default=128)
     parser.add argument('--epochs',
                             help='maximum number of iterations.'.
                             type=int,
                             default=50)
     parser.add_argument('--sample_size',
                             help='number of images to generate.',
                             type=int,
                             default=25)
     parser.add_argument('--latent-dim',
                             help='.',
                             type=int,
                             default=100)
     parser.add argument('--Ir',
                             help='initial learning rate.',
                             type=float,
                             default=1e-3)
     args = parser.parse args()
     main(args)
```

# Results

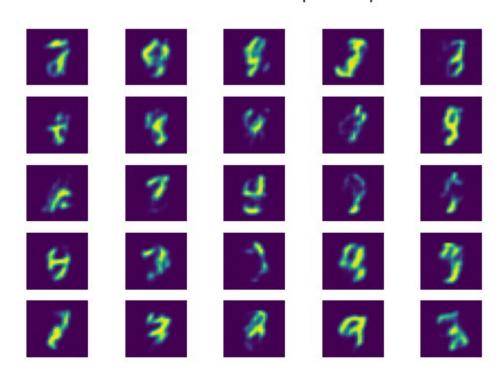
## 1. Dataset - MNIST

The following plot show the train and test ELBO values per epochs:



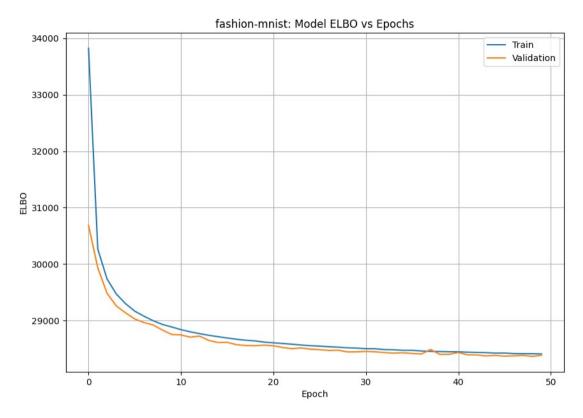
The following plot show 25 sampled images from the trained model:

mnist: Trained Model - Output Samples



### 2. Dataset - Fashion-MNIST

The following plot show the train and test ELBO values per epochs:



The following plot show 25 sampled images from the trained model:

fashion-mnist: Trained Model - Output Samples

